

# Kubernetes

## MAIN COMPONENTS (17.02.22)

### **Worker node**

simple server, physical or virtual machine

### **Container**

A container is a standard unit of software that packages up code and all its dependencies, so the application runs quickly and reliably from one computing environment to another.

### **Pod – abstraction of container**

The smallest unit of Kubernetes (usually 1 container per pod), creates a running environment (layer on top of container) for the purpose of abstraction away from container runtime/technology (elevate users away from deeper mechanics) so that users can replace them if required. Users only interact with the Kubernetes layer.

>Can be multiple containers within a pod in cases where there is a main application and helper container/side services which must run within the pod

### Pod IP

Each pod gets its own IP address meaning they can communicate to each other (private) E.g application pod can communicate to db pod using IP

>Pods are ephemeral so they can die easily – when this occurs a new pod is automatically created and is assigned a new IP – every time a pod dies we have to adjust the Ip for communication

### **Service - communication**

Service is a permanent IP address that can be attached to each pod – lifecycle of service and pod are not connected.

It also acts as a load balancer (sends to least busy pod)

### External Service

Allows container to be accessed from external source using an IP

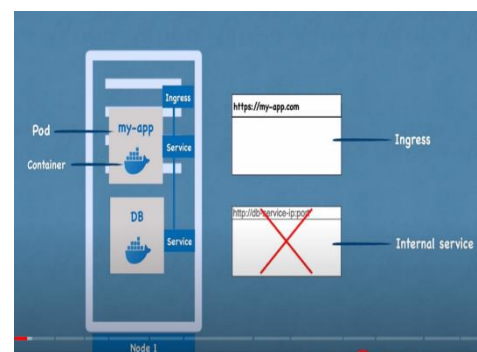

### Internal Service

Does not allow container to be accessed from an external source using an IP

### **Ingress – route traffic into cluster**

Service Ip often are unattractive consisting of service IP with the port number, Ingress offers a customizable IP which users access and it forwards them to the service IP.

### **ConfigMap – external configuration (plain text)**



The database URL is often in the built application (properties file) - if the URL changes you would have to adjust the URL in the application (rebuild, push to repo, pull to pod) – which is a lot of work.

ConfigMaps are used to manage the external configuration of your application (URL, External services etc). This is then connected to the pod – so when adjusting the name of a service we only have to update config map

### Secret – external configuration (encoded)

Secret operates like a config map but its used to store secret information such as user names and passwords (credentials) e.g user names, passwords etc. This is stored in base 64 encoded

>Can use data from config map or secrets using environment variables or as a properties file

### Volumes – data persistence

If the database or container gets restarted the data would be gone. Volumes attach a storage to your pod. This storage could be on a local machine (same server node that the pod is running) or remotely outside k8s cluster (e.g cloud or different server). The data is persisted and remains available after restarts.

>Have to recognise separation of k8s cluster and storage as k8s doesn't manage data persistence. You as a user or admin are responsible for backing up, replicating and ensuring the data is stored on appropriate hardware.

### Deployment – replicating mechanism (stateless)

If I restart my application or the pod dies, there's downtime where users cannot reach my application – this is bad In prod. Instead of relying on one node we replicate everything in a separate node and connect to the same service – as services act as a load balancer users are directed to the container which is least busy.

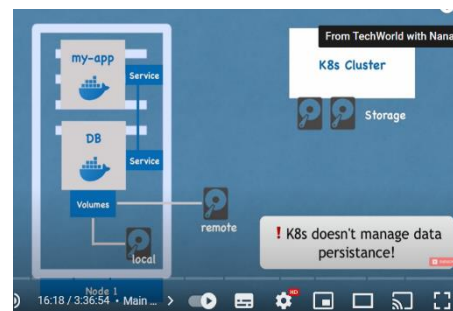
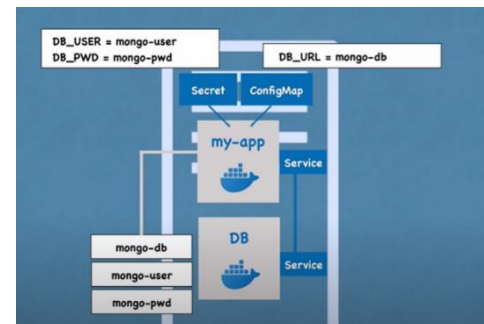
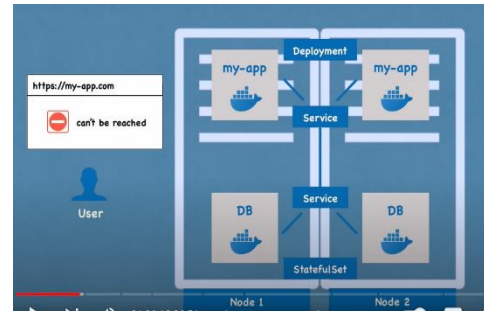
When creating the replica pod of our application we don't create a second pod and instead define a blueprint (deployment) and specify how many duplicates we require. In application we would not be creating pods we would instead be creating deployments – as we can specify how many instances we need and scale up and down based on requirements

>deployments are an abstraction of pods

>deployment for stateless Apps

### Statefulset– replicating mechanism (stateful)

Databases cannot be duplicated via deployment as they have states – if we have clones of DB they need to access shared data storage this means we need a mechanism which manages which pod can



write/read to/from the storage to avoid data inconsistencies. Stateful set is directed towards stateful apps – like deployment statefulset is responsible for replicating the pods and scaling them up or down

>statefulset for stateful apps or databases